

Screenshots showing both ASL gestures correctly identified:

- “Hi” correctly identified:

```
hi: 0.901553  
sup: 0.098447
```

- “Sup” correctly identified:

```
hi: 0.128932  
sup: 0.871068
```

Fixed software issue:

In section 6, the model is compiled using “`model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])`”, which is incorrect. This is because MSE is supposed to be used as a regression metric, rather than as a loss function. By replacing `loss='mse'` with `loss='categorical_crossentropy'` and `metrics=['mae']` with `metrics=['accuracy']`, we get the intended results. We know that the correct match is categorical crossentropy because our labels are one-hot encoded vectors (e.g. [1, 0] and [0, 1]). If we were using labels with integer class indices (0 or 1), we would use `sparse_categorical_crossentropy` instead.